

Momentum Contrast for Unsupervised Visual Representation Learning (MoCo)

Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, Ross Girshick
Facebook AI Research (FAIR)

Sungman, Cho.

Introduction

Introduction : Contrastive Loss

- **A common way of defining a loss function** is to measure the difference between a model's prediction and a fixed target.

Reconstructing : L1 / L2

Pre-defined categories : Cross-Entropy / Margin-based Loss

Introduction : Contrastive Loss

- A common way of defining a loss function is to measure the difference between a model's prediction and a fixed target.

Reconstructing : L1 / L2

Pre-defined categories : Cross-Entropy / Margin-based Loss

- **Contrastive losses** measure the **similarities of sample pairs** in a representation space.

Introduction : Contrastive Loss

- A common way of defining a loss function is to measure the difference between a model's prediction and a fixed target.


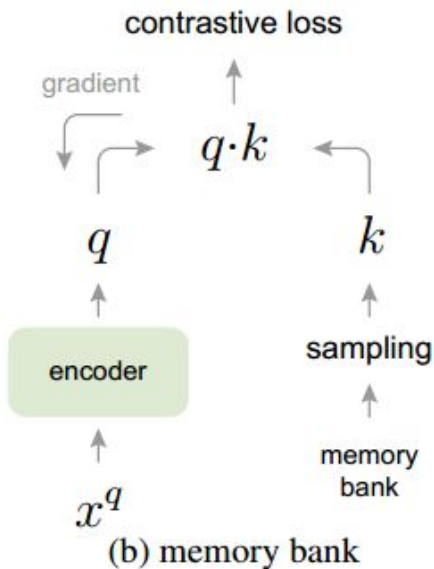

Reconstructing : L1 / L2

Pre-defined categories : Cross-Entropy / Margin-based Loss

- **Contrastive losses** measure the **similarities of sample pairs** in a representation space.
- **Adversarial losses** measure the **difference between probability distributions**.




Introduction

- Contrastive loss mechanisms



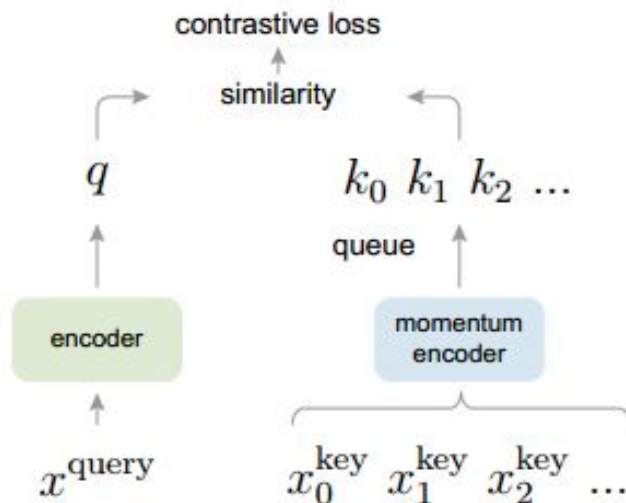
Introduction

- Contrastive loss mechanisms



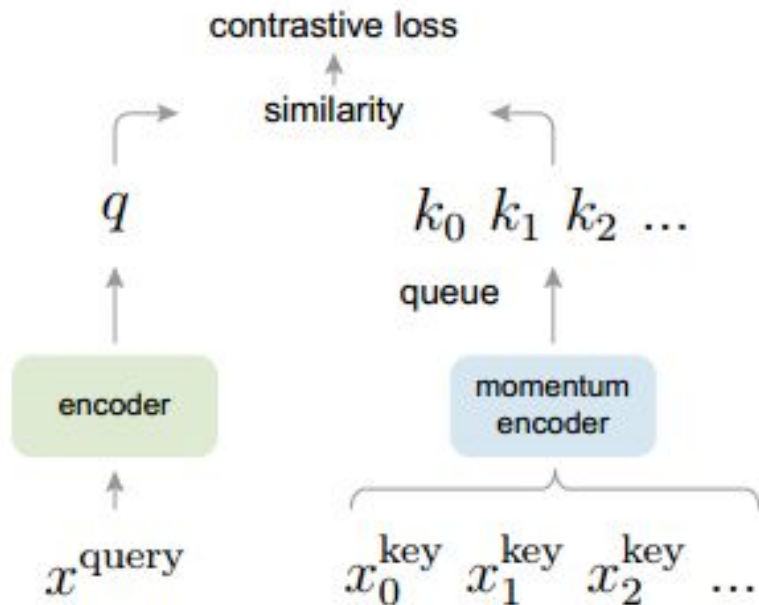
Contribution

- Build a **dynamic dictionary with a queue** and a moving-averaged encoder.
- Competitive results under the common linear protocol on ImageNet classification.



Method

Method : Contrastive as Dictionary Look-up



$$\mathcal{L}_q = -\log \frac{\exp(q \cdot \overset{\text{Matching key}}{k_+} / \tau)}{\sum_{i=0}^K \exp(q \cdot k_i / \tau)}$$

- The sum is over *one positive* and K *negative* samples

Method : Momentum Contrast

- Contrastive learning is a way of **building a discrete dictionary on high-dimensional continuous inputs** such as images.

Method : Momentum Contrast

- Contrastive learning is a way of **building a discrete dictionary on high-dimensional continuous inputs** such as images.
- Good features can be **learned by a large dictionary** that covers a **rich set of negative samples**, while the encoder for the **dictionary keys is kept as consistent** as possible despite its evolution.
- **Two subjects**
 - Dictionary as a queue.
 - Momentum update.

Dictionary as a queue

Method : Dictionary as a queue

- At the core of our approach is **maintaining the dictionary as a queue of data samples.**

Method : Dictionary as a queue

- At the core of our approach is maintaining the dictionary as a queue of data samples.
- The introduction of a **queue decouples the dictionary size from the mini-batch size.**

Method : Dictionary as a queue

- At the core of our approach is maintaining the dictionary as a queue of data samples.
- The introduction of a queue decouples the dictionary size from the mini-batch size.
- The samples in the dictionary are progressively replaced. The **current mini-batch is enqueued** to the dictionary, and the **oldest mini-batch in the queue is removed**.

Method : Dictionary as a queue

- At the core of our approach is maintaining the dictionary as a queue of data samples.
- The introduction of a queue decouples the dictionary size from the mini-batch size.
- The samples in the dictionary are progressively replaced. The current mini-batch is enqueued to the dictionary, and the oldest mini-batch in the queue is removed.
- Removing the **oldest mini-batch** can be beneficial, because its encoded keys are the most outdated and thus the **least consistent** with the newest ones.

Momentum update

Method : Momentum update.

- Using a queue can make **dictionary large**, but it also makes it **intractable to update** the key encoder by back-propagation (the gradient should propagate to all samples in the queue).
- **Naive solution** : copy the key encoder from the query encoder, ignoring this gradient.
- Propose a **momentum update** to address this issue.

Method : Momentum update.

Formally, denoting the parameters of f_k as θ_k and those of f_q as θ_q , we update θ_k by:

$$\theta_k \leftarrow m\theta_k + (1 - m)\theta_q. \quad (2)$$

Here $m \in [0, 1)$ is a momentum coefficient. Only the parameters θ_q are updated by back-propagation.

Pseudocode of MoCo

Algorithm 1 Pseudocode of MoCo in a PyTorch-like style.

```
# f_q, f_k: encoder networks for query and key
# queue: dictionary as a queue of K keys (CxK)
# m: momentum
# t: temperature

f_k.params = f_q.params # initialize
for x in loader: # load a minibatch x with N samples
    x_q = aug(x) # a randomly augmented version
    x_k = aug(x) # another randomly augmented version

    q = f_q.forward(x_q) # queries: NxC
    k = f_k.forward(x_k) # keys: NxK
    k = k.detach() # no gradient to keys

    # positive logits: Nx1
    l_pos = bmm(q.view(N,1,C), k.view(N,C,1))

    # negative logits: NxK
    l_neg = mm(q.view(N,C), queue.view(C,K))

    # logits: Nx(1+K)
    logits = cat([l_pos, l_neg], dim=1)

    # contrastive loss, Eqn. (1)
    labels = zeros(N) # positives are the 0-th
    loss = CrossEntropyLoss(logits/t, labels)

    # SGD update: query network
    loss.backward()
    update(f_q.params)




    # momentum update: key network
    f_k.params = m*f_k.params+(1-m)*f_q.params

    # update dictionary
    enqueue(queue, k) # enqueue the current minibatch
    dequeue(queue) # dequeue the earliest minibatch
```

bmm: batch matrix multiplication; mm: matrix multiplication; cat: concatenation.




Relations to previous mechanisms

Related Works



- Limited by the GPU memory size.

Related Works




- Can't update memory bank, So less consistent

Shuffling BN

- Using BN prevents the model from learning good representations.
- The model appears to “cheat” the pretext task and easily find a low-loss solution.
- This is possibly because the intra-batch communication among samples leaks information.
- Shuffle the sample order in the current mini-batch before distributing it among GPUs

Experiments

ImageNet



ImageNet

method	architecture	#params (M)	accuracy (%)
Exemplar [17]	R50w3 \times	211	46.0 [38]
RelativePosition [13]	R50w2 \times	94	51.4 [38]
Jigsaw [45]	R50w2 \times	94	44.6 [38]
Rotation [19]	Rv50w4 \times	86	55.4 [38]
Colorization [64]	R101*	28	39.6 [14]
DeepCluster [3]	VGG [53]	15	48.4 [4]
BigBiGAN [16]	R50	24	56.6
	Rv50w4 \times	86	61.3
<i>methods based on contrastive learning follow:</i>			
InstDisc [61]	R50	24	54.0
LocalAgg [66]	R50	24	58.8
CPC v1 [46]	R101*	28	48.7
CPC v2 [35]	R170* _{wider}	303	65.9
CMC [56]	R50 _{L+ab}	47	64.1 [†]
	R50w2 \times _{L+ab}	188	68.4 [†]
	AMDIM _{small}	194	63.5 [†]
AMDIM [2]	AMDIM _{large}	626	68.1 [†]
MoCo	R50	24	60.6
	RX50	46	63.9
	R50w2 \times	94	65.4
	R50w4 \times	375	68.6

Momentum

momentum m	0	0.9	0.99	0.999	0.9999
accuracy (%)	<i>fail</i>	55.2	57.8	59.0	58.9

Shuffling BN




Figure A.1. **Ablation of Shuffling BN.** *Dash:* training curve of the pretext task, plotted as the accuracy of $(K+1)$ -way dictionary lookup. *Solid:* validation curve of a kNN-based monitor [61] (not a linear classifier) on ImageNet classification accuracy. This plot shows the first 80 epochs of training: training longer without shuffling BN overfits more.

Thank you.