

# **ALBERT: A LITE BERT FOR SELF-SUPERVISED LEARNING OF LANGUAGE REPRESENTATIONS**

# word2vec

<https://ratsgo.github.io/from%20frequency%20to%20semantics/2017/03/30/word2vec/>

<https://mc.ai/deep-nlp-word-vectors-with-word2vec/>

<https://wikidocs.net/22660>

# Bert

<http://jalammar.github.io/illustrated-transformer/>

<https://medium.com/dissecting-bert/dissecting-bert-part-1-d3c3d495cdb3>

# 1. BackGround(NLP, Text Mining)

## 1. Text embedding(encoding, representation vector)

- a. How Text can be represented to vector(number) for feeding to model

## 2. Text tokenization

- a. rule 기반(ex konlp기반 형태소 분석기)
- b. wordpiece 기반( ex sentence2vec)
- c. 나는 오늘 식당에 가서 밥을 먹었다 -> 나 는 오늘 식당 에 가 서 밥 을 먹었 다

## 3. one hot vector(sparse vector)

- a. 나는 오늘 식당에서 밥을 먹었다. -> 1 2 3 4 5  
그리고 나는 식당에서 책도 읽었다. -> 6 1 3 7 8

# 1. Background(NLP, Text Mining)

## 1. Bag of words

a. 나는(1) 오늘(2) 식당에서(3) 밥을(4) 먹었다(5) 그리고(6) 책도(7) 읽었다(8)

b. 1 1 1 1 1 0 0 0

c. 1 0 1 0 0 1 1 1

## 2. tfidf

a. 단어별 중요도에 따른 가중치 추가

b. 0.5 0.3 1.5 1.2 1 0 0 0

## 3.

# 1. BackGround(NLP, Text Mining)

신경망을 이용한 text embedding (word2vec, fasttext, glove (using NN))

중심 단어  
↓  
주변 단어

The fat cat sat on the mat

The fat cat sat on the mat

The fat cat sat on the mat


The fat cat sat on the mat

The fat cat sat on the mat



The fat cat sat on the mat

The fat cat sat on the mat

중심 단어	주변 단어
[1, 0, 0, 0, 0, 0, 0]	[0, 1, 0, 0, 0, 0, 0], [0, 0, 1, 0, 0, 0, 0]
[0, 1, 0, 0, 0, 0, 0]	[1, 0, 0, 0, 0, 0, 0], [0, 0, 1, 0, 0, 0, 0], [0, 0, 0, 1, 0, 0, 0]
[0, 0, 1, 0, 0, 0, 0]	[1, 0, 0, 0, 0, 0, 0], [0, 1, 0, 0, 0, 0, 0], [0, 0, 0, 1, 0, 0, 0], [0, 0, 0, 0, 1, 0, 0]
[0, 0, 0, 1, 0, 0, 0]	[0, 1, 0, 0, 0, 0, 0], [0, 0, 1, 0, 0, 0, 0], [0, 0, 0, 0, 1, 0, 0], [0, 0, 0, 0, 0, 1, 0]
[0, 0, 0, 0, 1, 0, 0]	[0, 0, 1, 0, 0, 0, 0], [0, 0, 0, 1, 0, 0, 0], [0, 0, 0, 0, 0, 1, 0], [0, 0, 0, 0, 0, 0, 1]
[0, 0, 0, 0, 0, 1, 0]	[0, 0, 0, 1, 0, 0, 0], [0, 0, 0, 0, 1, 0, 0], [0, 0, 0, 0, 0, 0, 1]
[0, 0, 0, 0, 0, 0, 1]	[0, 0, 0, 0, 1, 0, 0], [0, 0, 0, 0, 0, 1, 0]



# 1. BackGround(NLP, Text Mining)



# 1. BackGround(NLP, Text Mining)

## 1. 언어 모델링

- 이전 단어들이 주어졌을 때 다음 단어를 예측
- 통계를 이용한 방법과 인공 신경망을 이용한 방법
- Ngram vs 신경망(LSTM)

n-gram을 통한 언어 모델에서는 다음에 나올 단어의 예측은 오직 n-1개의 단어에만 의존합니다. 예를 들어 'An adorable little boy is spreading' 다음에 나올 단어를 예측하고 싶다고 할 때, n=4라고 한 4-gram을 이용한 언어 모델을 사용한다고 합시다. 이 경우, spreading 다음에 올 단어를 예측하는 것은 n-1에 해당되는 앞의 3개의 단어만을 고려합니다.

~~An adorable little~~ boy is spreading ?  
무시됨!      n-1개의 단어

$$P(w|\text{boy is spreading}) = \frac{\text{count}(\text{boy is spreading } w)}{\text{count}(\text{boy is spreading})}$$

만약 갖고있는 코퍼스에서 boy is spreading가 1,000번 등장했다고 합시다. 그리고 boy is spreading insults가 500번 등장했으며, boy is spreading smiles가 200번 등장했다고 합시다. 그렇게 되면 boy is spreading 다음에 insults가 등장할 확률은 50%이며, smiles가 등장할 확률은 20%입니다. 확률적 선택에 따라 우리는 insults가 더 맞다고 판단하게 됩니다.

$$P(\text{insults}|\text{boy is spreading}) = 0.500$$

$$P(\text{smiles}|\text{boy is spreading}) = 0.200$$




# 1. BackGround(NLP, Text Mining)

## 1. 사전학습 (word2vec etc)

- Use more Deep NN -> ELMO, OPEN-GPT, BERT
- unidirectional vs bidirectional ( Open GPT vs ELMO, BERT)
- feature based vs fine tuning ( ELMO vs OPEN GPT, BERT)

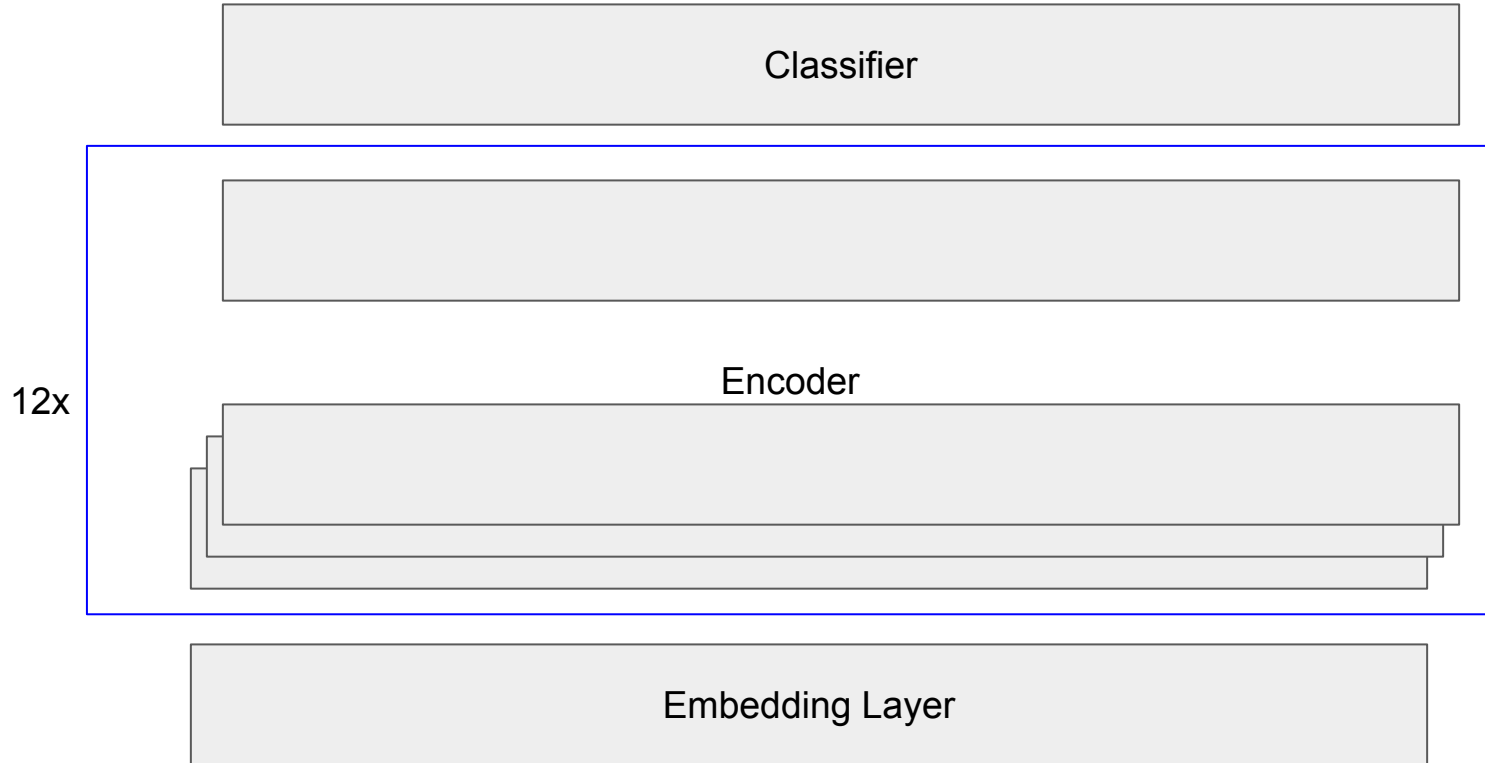
## 2. ALBERT

- BERT is very Good, but too much computing resource, long training time
- minimized the parameter with maintaining performance




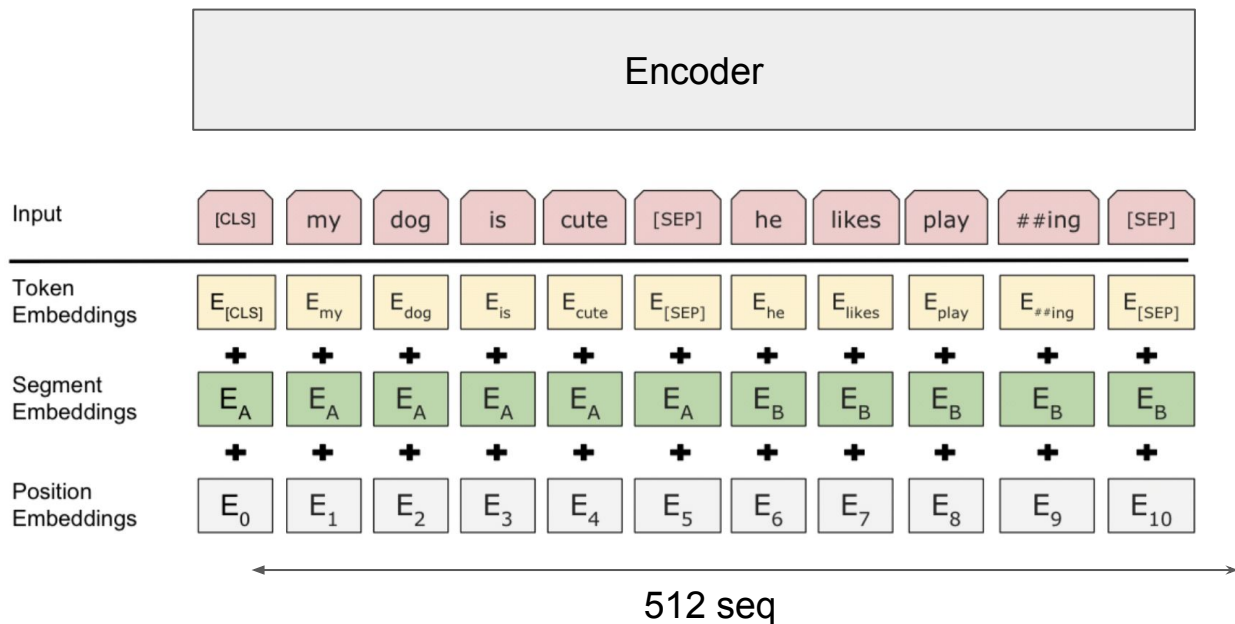


# 1. BERT(Bidirectional Encoder representation of Transformer)



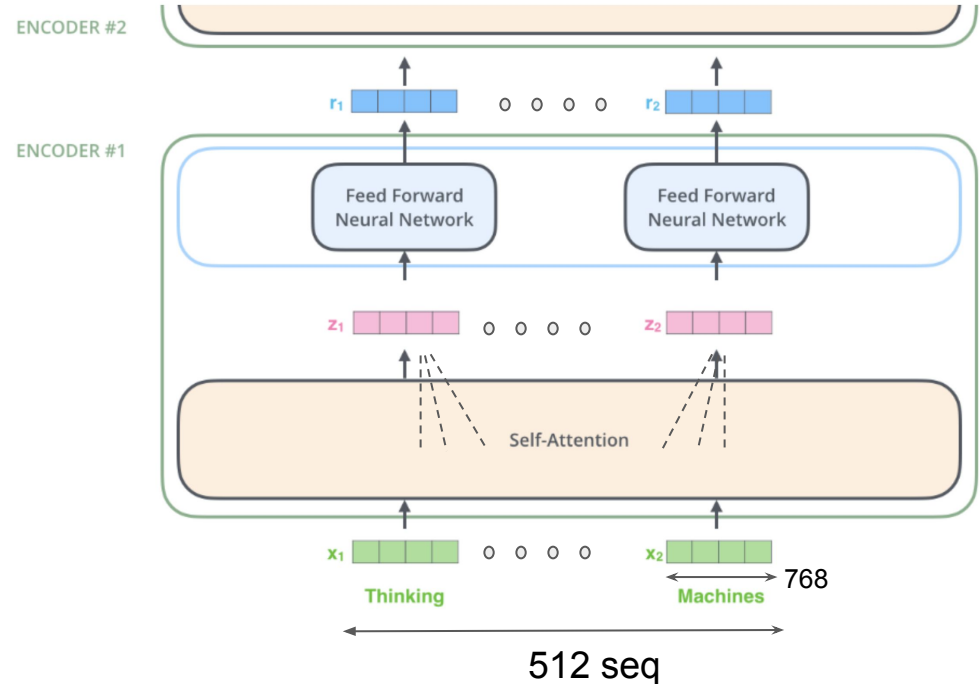

# Embedding, positional encoding

give offset according to words position



# 1. BERT(Bidirectional Encoder representation of Transformer)

It similar to FCN (not scalar but vector, how to calculate similarity ? use dot product)





Input




Thinking

Machines

Embedding





Queries






$W^Q$

Keys



$W^K$

Values



$W^V$

Input

Embedding

Queries

Keys

Values

Score

Divide by 8 (  $\sqrt{d_k}$  )

Softmax

Softmax


X

Value


Sum

Thinking


$x_1$




$q_1$



$k_1$



$v_1$




$$q_1 \bullet k_1 = 112$$


14

0.88

$v_1$




$z_1$




Machines


$x_2$




$q_2$



$k_2$



$v_2$




$$q_1 \bullet k_2 = 96$$


12

0.12


$v_2$



$z_2$



how about offset or activation ft?



similar to convolution in image  
1d convolution

$$\text{softmax} \left( \frac{Q \times K^T}{\sqrt{d_k}} \right) V$$

$Z$

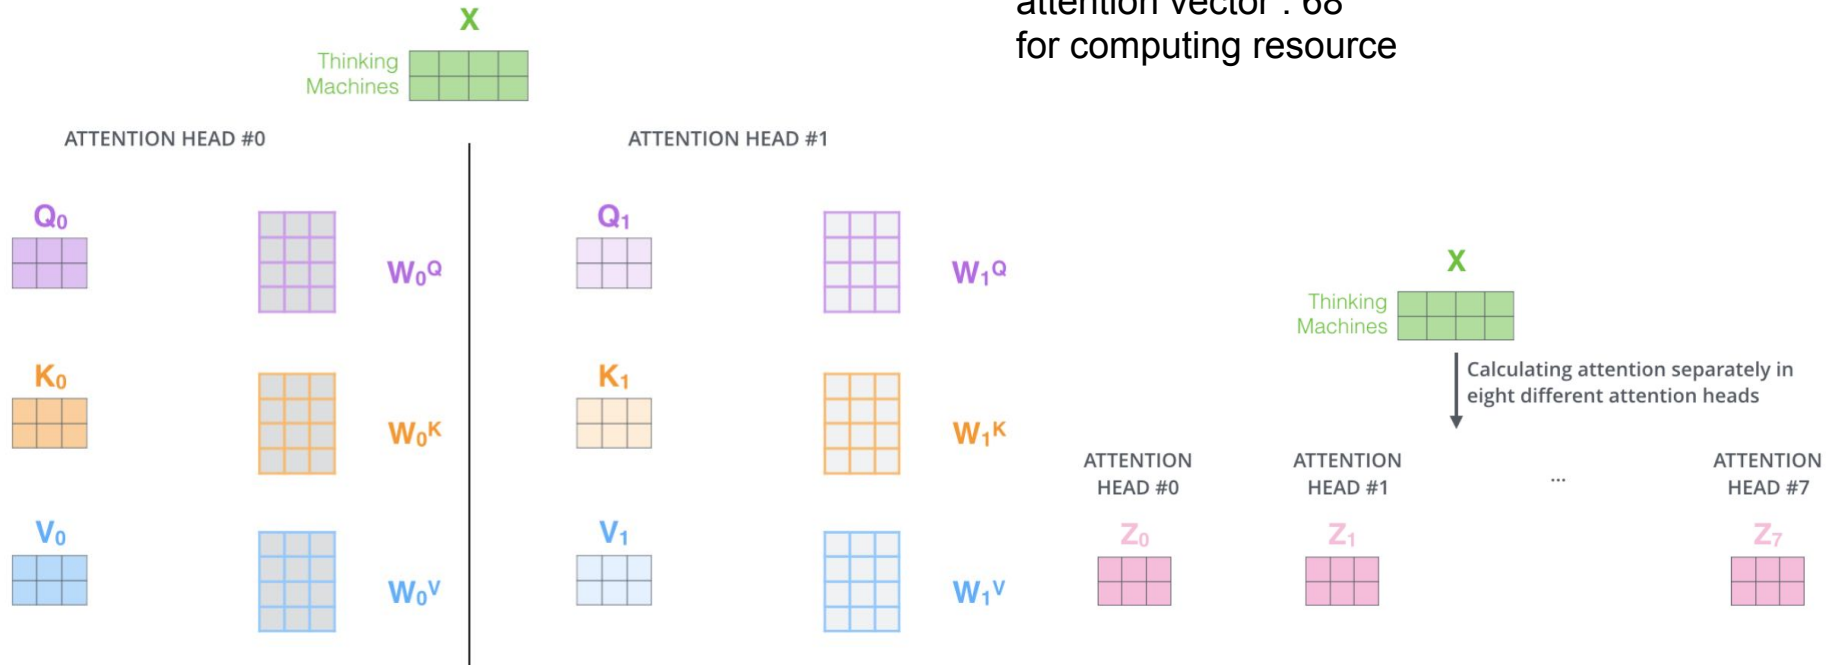
multi head attention

768  $\rightarrow$  68 \* 8

embedding vector : 768

attention vector : 68

for computing resource




1) Concatenate all the attention heads



2) Multiply with a weight matrix  $W^O$  that was trained jointly with the model

X



3) The result would be the  $Z$  matrix that captures information from all the attention heads. We can send this forward to the FFNN





1) This is our input sentence\*

2) We embed each word\*

3) Split into 8 heads.  
We multiply  $X$  or  $R$  with weight matrices




4) Calculate attention using the resulting  $Q/K/V$  matrices

5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer

Thinking  
Machines





\* In all encoders other than #0, we don't need embedding.  
We start directly with the output of the encoder right below this one

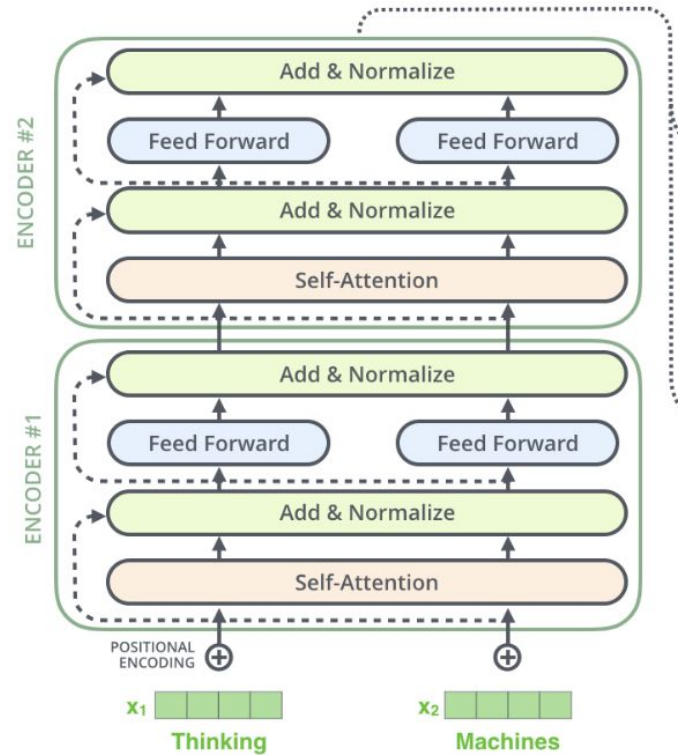
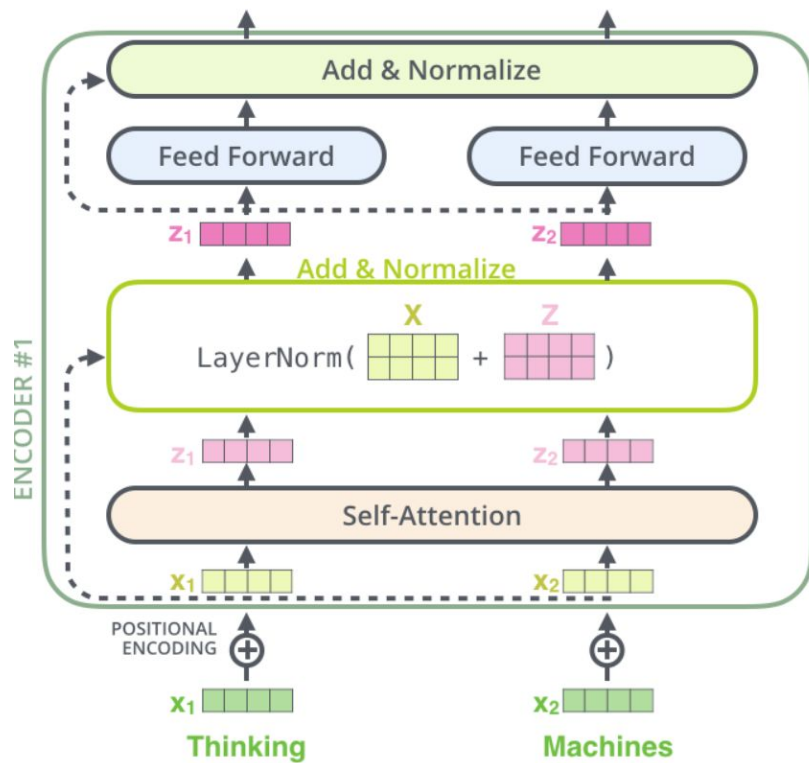


...

...

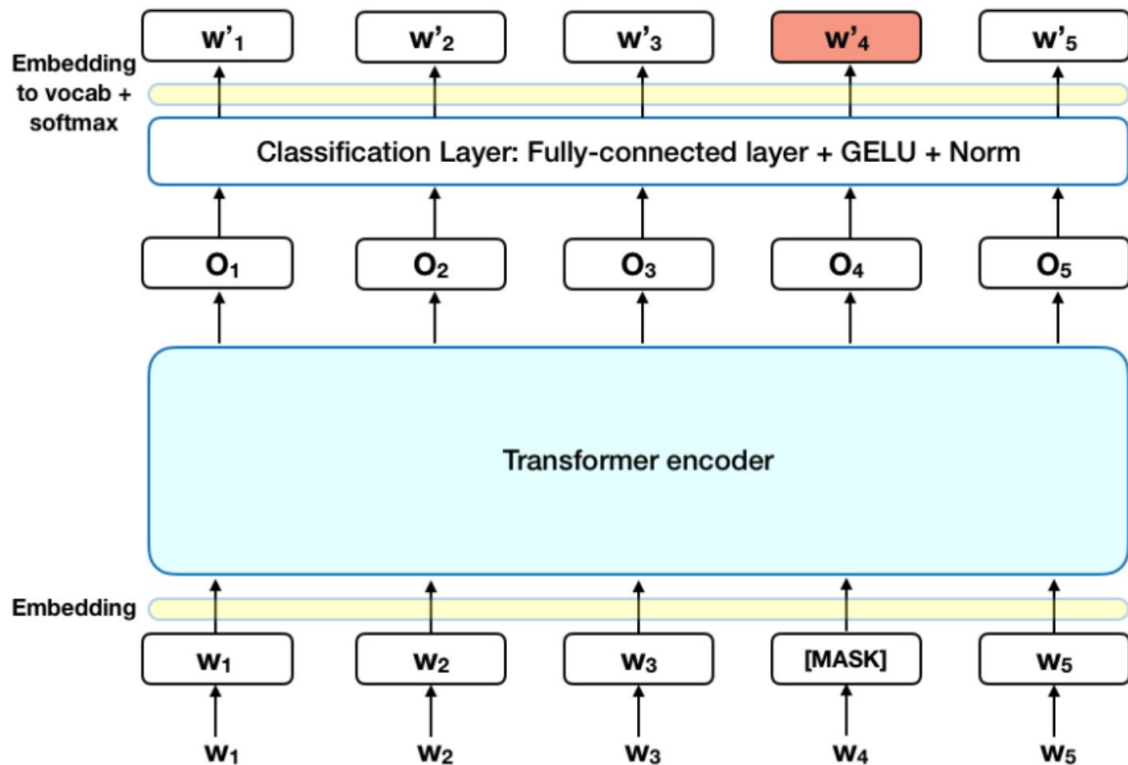
...





# How to train bert?

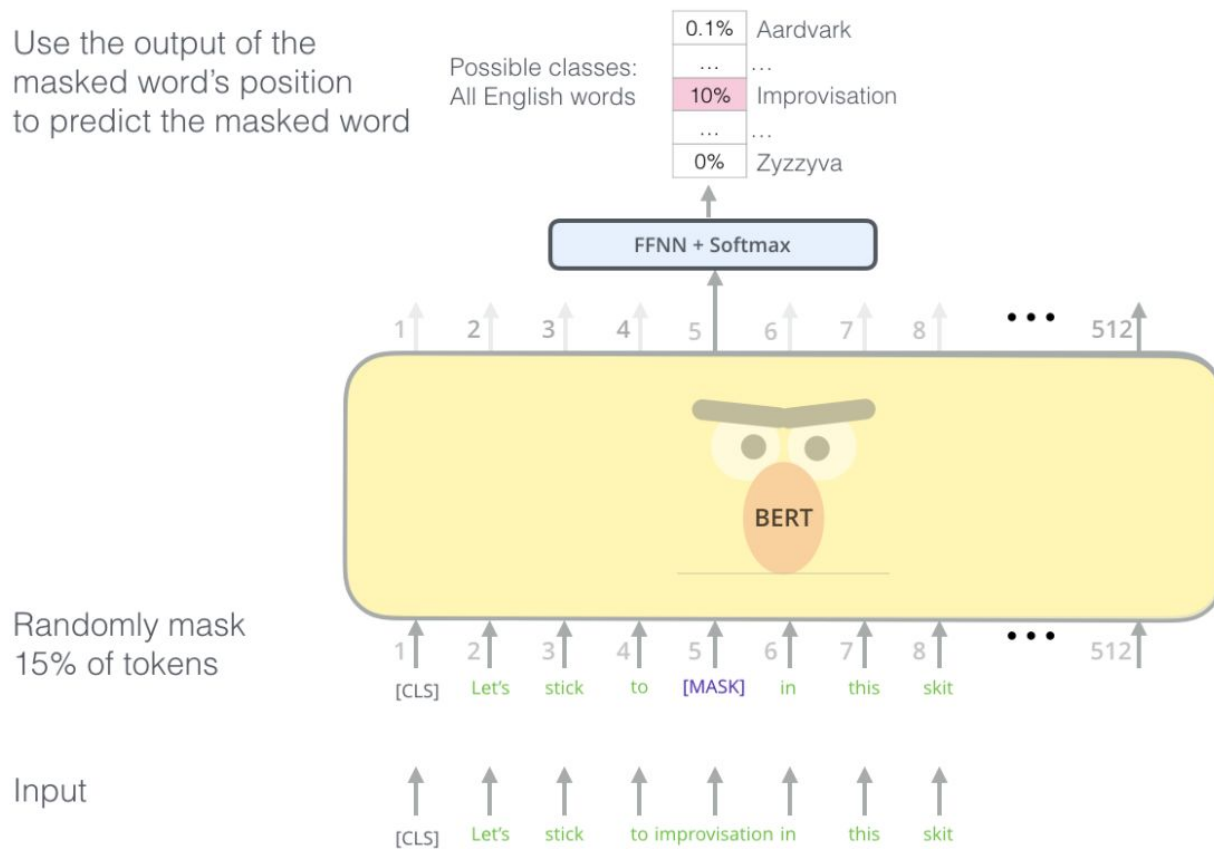
## 1. MLM(masked Language model)



masking rate : 15%  
masking : 80, wrong, 10, correct 10

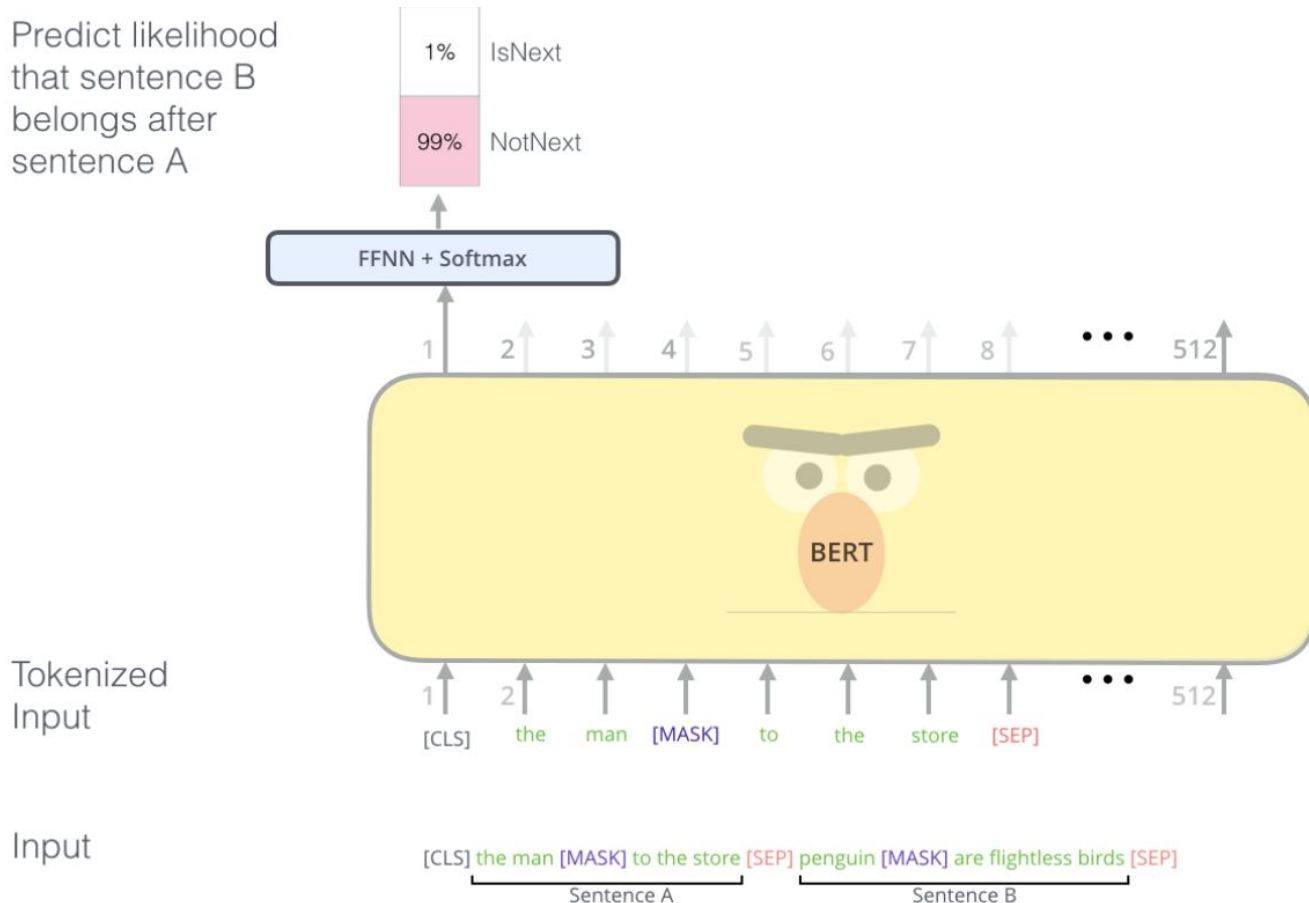
# Masked language model

Use the output of the masked word's position to predict the masked word



# NSP(next sentence prediction)

Predict likelihood  
that sentence B  
belongs after  
sentence A



# ALBERT

Factorized embedding parameterization

Parameter sharing

NSO

Model		Parameters	Layers	Hidden	Embedding	Parameter-sharing
BERT	base	108M	12	768	768	False
	large	334M	24	1024	1024	False
ALBERT	base	12M	12	768	128	True
	large	18M	24	1024	128	True
	xlarge	60M	24	2048	128	True
	xxlarge	235M	12	4096	128	True

Model		Parameters	SQuAD1.1	SQuAD2.0	MNLI	SST-2	RACE	Avg	Speedup
BERT	base	108M	90.4/83.2	80.4/77.6	84.5	92.8	68.2	82.3	4.7x
	large	334M	92.2/85.5	85.0/82.2	86.6	93.0	73.9	85.2	1.0
ALBERT	base	12M	89.3/82.3	80.0/77.1	81.6	90.3	64.0	80.1	5.6x
	large	18M	90.6/83.9	82.3/79.4	83.5	91.7	68.5	82.4	1.7x
	xlarge	60M	92.5/86.1	86.1/83.1	86.4	92.4	74.8	85.5	0.6x
	xxlarge	235M	<b>94.1/88.3</b>	<b>88.1/85.1</b>	<b>88.0</b>	<b>95.2</b>	<b>82.3</b>	<b>88.7</b>	0.3x

# ALBERT

Factorized embedding parameterization

bert : embedding == hidden size(768)

albert : embedding =128, hidden size :768

1. wordpiece embedding meant to learn context independent representation, hidden size embedding meant to learn context dependent representation
2. Parameter decrease :  $30000(\text{vocab size}) * 768 \rightarrow 30000 * 128 + 128 * 768$

Model	<i>E</i>	Parameters	SQuAD1.1	SQuAD2.0	MNLI	SST-2	RACE	Avg
ALBERT base not-shared	64	87M	89.9/82.9	80.1/77.8	82.9	91.5	66.7	81.3
	128	89M	89.9/82.8	80.3/77.3	83.7	91.5	67.9	81.7
	256	93M	90.2/83.2	80.3/77.4	84.1	91.9	67.3	81.8
	768	108M	90.4/83.2	80.4/77.6	84.5	92.8	68.2	82.3
ALBERT base all-shared	64	10M	88.7/81.4	77.5/74.8	80.8	89.4	63.5	79.0
	128	12M	89.3/82.3	80.0/77.1	81.6	90.3	64.0	80.1
	256	16M	88.8/81.5	79.1/76.3	81.5	90.3	63.4	79.6
	768	31M	88.6/81.5	79.2/76.6	82.0	90.6	63.3	79.8

# ALBERT

## CROSS-LAYER PARAMETER SHARING

group sharing(for Layer)

	Model	Parameters	SQuAD1.1	SQuAD2.0	MNLI	SST-2	RACE	Avg
ALBERT base $E=768$	all-shared	31M	88.6/81.5	79.2/76.6	82.0	90.6	63.3	79.8
	shared-attention	83M	89.9/82.7	80.0/77.2	84.0	91.4	67.7	81.6
	shared-FFN	57M	89.2/82.1	78.2/75.4	81.5	90.8	62.6	79.5
	not-shared	108M	90.4/83.2	80.4/77.6	84.5	92.8	68.2	82.3
ALBERT base $E=128$	all-shared	12M	89.3/82.3	80.0/77.1	82.0	90.3	64.0	80.1
	shared-attention	64M	89.9/82.8	80.7/77.9	83.4	91.9	67.6	81.7
	shared-FFN	38M	88.9/81.6	78.6/75.6	82.3	91.7	64.4	80.2
	not-shared	89M	89.9/82.8	80.3/77.3	83.2	91.5	67.9	81.6



# ALBERT

## SENTENCE ORDER PREDICTION (SOP)

SP tasks	Intrinsic Tasks			Downstream Tasks					
	MLM	NSP	SOP	SQuAD1.1	SQuAD2.0	MNLI	SST-2	RACE	Avg
None	54.9	52.4	53.3	88.6/81.5	78.1/75.3	81.5	89.9	61.7	79.0
NSP	54.5	90.5	52.0	88.4/81.5	77.2/74.6	81.6	<b>91.1</b>	62.3	79.2
SOP	54.0	78.9	86.5	<b>89.3/82.3</b>	<b>80.0/77.1</b>	<b>82.0</b>	90.3	<b>64.0</b>	<b>80.1</b>

# ALBERT

Models	MNLI	QNLI	QQP	RTE	SST	MRPC	CoLA	STS	WNLI	Avg
<i>Single-task single models on dev</i>										
BERT-large	86.6	92.3	91.3	70.4	93.2	88.0	60.6	90.0	-	-
XLNet-large	89.8	93.9	91.8	83.8	95.6	89.2	63.6	91.8	-	-
RoBERTa-large	90.2	94.7	<b>92.2</b>	86.6	96.4	<b>90.9</b>	68.0	92.4	-	-
ALBERT (1M)	90.4	95.2	92.0	88.1	96.8	90.2	68.7	92.7	-	-
ALBERT (1.5M)	<b>90.8</b>	<b>95.3</b>	<b>92.2</b>	<b>89.2</b>	<b>96.9</b>	<b>90.9</b>	<b>71.4</b>	<b>93.0</b>	-	-
<i>Ensembles on test (from leaderboard as of Sept. 16, 2019)</i>										
ALICE	88.2	95.7	<b>90.7</b>	83.5	95.2	92.6	<b>69.2</b>	91.1	80.8	87.0
MT-DNN	87.9	96.0	89.9	86.3	96.5	92.7	68.4	91.1	89.0	87.6
XLNet	90.2	98.6	90.3	86.3	96.8	93.0	67.8	91.6	90.4	88.4
RoBERTa	90.8	98.9	90.2	88.2	96.7	92.3	67.8	92.2	89.0	88.5
Adv-RoBERTa	91.1	98.8	90.3	88.7	96.8	93.1	68.0	92.4	89.0	88.8
ALBERT	<b>91.3</b>	<b>99.2</b>	90.5	<b>89.2</b>	<b>97.1</b>	<b>93.4</b>	69.1	<b>92.5</b>	<b>91.8</b>	<b>89.4</b>